

CSE 333 Notes

Contents

1	Memory & Data	3
1.1	Pointers	3
1.2	Pointer Arithmetic	3
1.3	Arrays	3
1.4	Array Subscript Notation	3
1.5	Strings	3
1.6	String Literals	3
1.7	Memory Layout	4
1.8	Significant Bits	4
1.9	Endianness	4
2	Programming in C	5
2.1	Generic Program Layout	5
2.2	Functions	5
2.3	Error Handling	5
2.4	Status Codes and Signals	6
2.5	Modules	6
2.6	Header Files	6
2.7	Header Guards	6
2.8	Preprocessors	6
2.9	Linkages	7
2.10	Format Specifiers	7
3	Memory Allocation	8
3.1	Null	8
3.2	Dynamic Memory Allocation	8
3.3	Dynamic Memory Deallocation	8
3.4	Structs	9
3.5	Typedef	9
3.6	Use of Structs	9
4	File I/O	10
4.1	Streams	10
4.2	Stream Functions	10
4.3	Stream Error Checking & Handling	10
4.4	Stream Buffering	11
4.5	Portable Operating System Interface (POSIX)	11
4.6	Directories	11
5	Programming in C++	12
5.1	Pointers	12
5.2	References	12
5.3	const	12
5.4	Classes	12
5.5	Class Definitions	13
5.6	Class Declarations	13
5.7	Constructors	13
5.8	Destructors	13

5.9	Rule of Three	13
5.10	Access Modifiers	13
5.11	Non-Member Functions	14
5.12	<code>friend</code> Non-Member Functions	14
5.13	Namespaces	14
5.14	<code>new/delete</code>	14
5.15	Templates	14
6	C++ Standard Template Library	15
6.1	Type Inference	15
6.2	Containers	15
6.3	Lists	15
6.4	Vectors	15
6.5	Maps	15
6.6	Iterators	16
7	Smart Pointers	17
7.1	Smart Pointers	17
7.2	<code>shared_ptr</code>	17
7.3	<code>unique_ptr</code>	17
7.4	<code>weak_ptr</code>	17
8	Inheritance	18
8.1	Inheritance	18
8.2	Virtual Functions	18
8.3	Overriding Functions	18
8.4	Pure Virtual Functions	18
8.5	Abstract Classes	18
9	Casts	19
9.1	Static Casts	19
9.2	Dynamic Casts	19
9.3	Constant Casts	19
9.4	Reinterpret Casts	19

1 Memory & Data

1.1 Pointers

Pointers are special variables that store addresses

- Since the length of an address is the word size, the size of a pointer is also the word size
- The address-of operator (&) returns the address of a variable in memory
- The dereference operator (*) accesses the data at the location of the pointer
- `type*` denotes the address of a variable of type `type`

1.2 Pointer Arithmetic

Pointer arithmetic takes the size of the data type being pointed to into account by automatically scaling the arithmetic operation

- i.e. Given `int* p`, the operation `p += 1` increases the value of `p` by 4, the size of `int`
- i.e. Given `long* p`, the operation `p += 1` increases the value of `p` by 8, the size of `long`
- i.e. Given `int* p1` and `int* p2`, the operation `p2 - p1` returns the number of `int` values that fit between the two addresses

1.3 Arrays

An array is a set of contiguous locations in memory that store the same type of data object

- Declaring an array follows the syntax `type array_name[num]`; and sets aside `num * sizeof(type)` bytes of consecutive memory
 - `num` is the number of elements in the array

1.4 Array Subscript Notation

`array_name[n]` refers to the n^{th} element of the array, and is actually accessed via `*(array_name + n)`

- In the C programming language, `n` can be smaller than 0 or larger than `num`
 - This will access an address beyond the scope of the array

1.5 Strings

A string is represented by an array of characters that are terminated by the null character

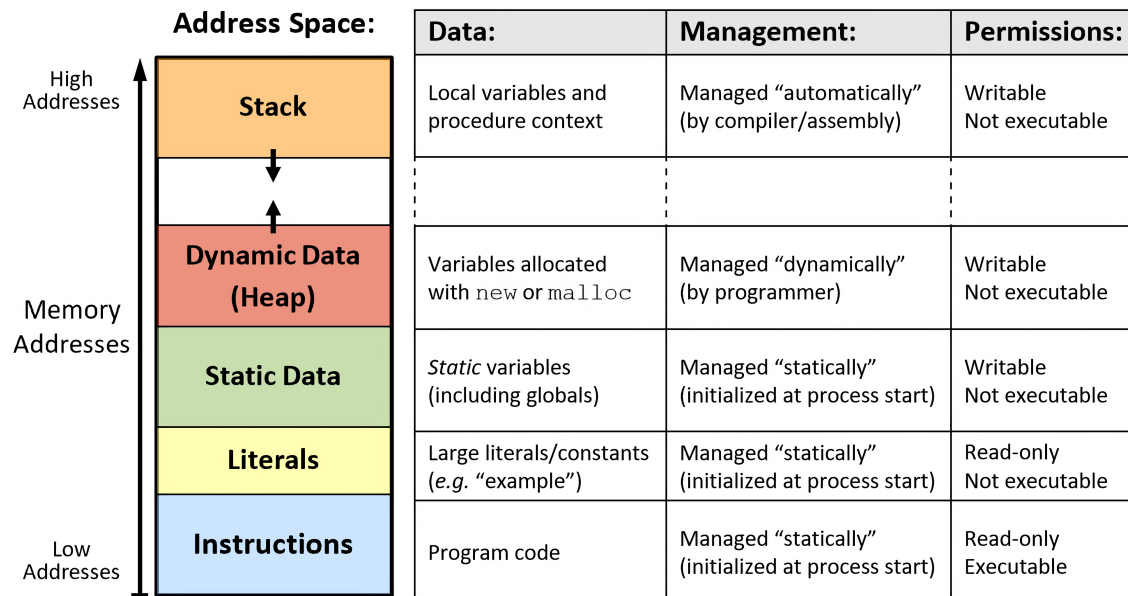
- i.e. `char str[] = "hello"` will allocate a character array of size 6 with the following content

'h'	'e'	'l'	'l'	'o'	'\0'
-----	-----	-----	-----	-----	------

1.6 String Literals

A string literal is a sequence of characters surrounded by double quotes (i.e. `"hello world"`) and automatically stored in memory as an array of characters, terminated by `'\0'`

1.7 Memory Layout



1.8 Significant Bits

- The left-most bit in a chunk of memory is the most-significant bit
- The right-most bit in a chunk of memory is the least-significant bit

1.9 Endianness

- In a big-endian machine, the most significant byte of data is stored at the lowest address
- In a little-endian machine, the least significant byte of data is stored at the lowest address

2 Programming in C

2.1 Generic Program Layout

```

1  #include <system_files>
2  #include "local_files"
3
4  #define macro_name macro_expr
5
6  /* declare functions */
7  /* declare external variables & structs */
8
9  int main(int argc, char* argv[]) {
10     /* statements */
11 }
12
13 /* define other functions */

```

- `argc` is the number of command-line arguments passed
- `argv` is an array containing *pointers* to the command-line arguments which are represented as strings
- The executable name is passed in as an argument

2.2 Functions

```

1  /* declare function */
2  returnType fname(type param1, ..., type paramN);
3
4  /* call function */
5  int main(int argc, char* argv[]) {
6     fname(arg1, ..., argN);
7  }
8
9  /* define function */
10 returnType fname(type param1, ..., type paramN) {
11     /* statements */
12 }

```

- Functions must be declared before being called
 - Functions are often declared in header files and incorporated via `#include`
- Functions can only be defined once
 - Functions cannot be overloaded

2.3 Error Handling

- C does not have error handling, i.e. `try...catch` block
- Functions return errors as integer error codes
 - `CONSTANT_NAMES` are used to abstract away the integer error codes
- The global variable `errno` holds the value of the last system error

2.4 Status Codes and Signals

- Processes exit with status codes, e.g. `EXIT_SUCCESS` and `EXIT_FAILURE`
 - Standard codes can be found in `stdlib.h`
- Crashes trigger signals from the OS

2.5 Modules

A module is a self-contained piece of an overall program

- Has externally visible functions that can be invoked
- Has externally visible `typedefs` and global variables that can be used
- May have internal functions, `typedefs`, and global variables that are private
- Has an interface which declares the module's public functions, `typedefs`, and global variables

2.6 Header Files

A header file is a file which declares the interface to a module

- The header `main.h` corresponds to the source file `main.c`
- Holds the variables, types, and function prototype declarations that make up the interface to a module
- Enables the use of `#include` to import the module
- Documentation must be written in the header file

2.7 Header Guards

Header guards safeguard against redefining macros

```

1  #ifndef PAIR_H_
2
3  struct Pair {
4      int a;
5      int b;
6  }
7  #endif

```

- Use header guards to define macros everywhere they are used

2.8 Preprocessors

The C preprocessor is a sequential and stateful search-and-replace text-processor that transforms your source code before the compiler runs

- The preprocessor takes in a C file and outputs a C file
- The preprocessor processes the directives found in the file, i.e. `#include`, `#define`, `//comments`

2.9 Linkages

Linkages specify whether an identifier declared in different scopes or in the same scope more than once refer to the same object or function

- External linkage
 - `extern` makes a declaration externally visible
 - i.e. `extern int i;`
 - Global variables and functions are `extern` by default
- Internal linkage
 - `static` restricts a definition to visibility within the file
 - i.e. `static int i = 0;`
 - It is good practice to use `static` when declaring local globals

2.10 Format Specifiers

Specifier	Output
<code>%d</code>	Signed int
<code>%u</code>	Unsigned int
<code>%o</code>	Unsigned octal
<code>%x</code>	Unsigned hex int (lower)
<code>%X</code>	Unsigned hex int (upper)
<code>%f</code>	Decimal floating point
<code>%a</code>	Hexadecimal floating point (lower)
<code>%A</code>	Hexadecimal floating point (upper)
<code>%c</code>	Character
<code>%s</code>	String
<code>%p</code>	Pointer
<code>%n</code>	Return index of pointer

- Use the `0N` modifier to print an int with at least `N` digits
 - i.e. `printf("%03d", 3)` outputs `003`
- Use the `0.N` modifier to print a float with exactly `N` decimal places
 - i.e. `printf("%03f", 3.14)` outputs `3.140`
- Use the `0N.M` modifier to print a float with at least `N – 1` digits and exactly `N` decimal places
 - i.e. `printf("%07.4f", 3.14)` outputs `03.1400`

3 Memory Allocation

3.1 Null

`NULL` is a memory location that is guaranteed to be invalid

- Dereferencing `NULL` will cause a segmentation fault
- `NULL` is often used as an indicator of an uninitialized pointer or an allocation error

3.2 Dynamic Memory Allocation

`void* malloc(size_t size)` will request a continuous block of at least `size` bytes of uninitialized memory from the allocator

- `malloc` returns a pointer to the beginning of the allocated space, or `NULL` in case of an allocation failure
- Often `size` is calculated by use of the `sizeof()` macro
 - i.e. `int* ptr = (int*) malloc(n * sizeof(int))` allocates the space for an array of `n` `int`'s
- `calloc` is an allocation function that initializes the allocated heap block to all zeros
- `realloc` changes the size (up or down) of a previously allocated block

3.3 Dynamic Memory Deallocation

`void free(void* p)` will deallocate the entirety of the heap block pointed to by `p`

- `p` must hold the same address that was returned by the allocating function, `p` cannot point to an arbitrary byte within the heap block
- If `p` is `NULL`, then nothing happens
- Invalid calls to `free` will result in a system exception
 - i.e. address of the middle of a block, address of a previously deallocated block, address outside of the heap

3.4 Structs

A struct is a user-defined, structured group of variables. A struct definition is formatted as

```

1  struct struct_tag {
2      type_1 field_1;
3      ...
4      type_N field_N;
5  };
6  struct struct_tag var_name;
```

- The user-chosen struct tag is part of the name of the new data type we are defining, which is the two-word name `struct struct_tag`
 - The struct tag can be omitted if the struct is not used elsewhere in the code
- Fields can be accessed from an instance using the `'.'` operator, i.e. `struct.x`
- Fields can be accessed from a pointer
 - By dereferencing and using the `'.'` operator, i.e. `(*ptr).x`
 - By using the `->` operator, i.e. `ptr->x`

3.5 Typedef

Typedef allows us to create aliases to other data types, such as two-word struct data type names

- Typedef statements are of the form `typedef <data type> <alias>;`

i.e. for structs, a typedef statement can be used after or combined with the struct definition

```

1  struct point_st {
2      int x;
3      int y;
4  };
5  typedef struct point_st Point;
6  Point pt1;
```

```

1  typedef struct point_st {
2      int x;
3      int y;
4  } Point;
5  Point pt1;
```

3.6 Use of Structs

- Structs can be copied by assigning an instance of a struct to another instance
- If the struct is small and we are only reading it, then passing a copy of the struct is faster. Otherwise use pointers

4 File I/O

4.1 Streams

A stream is a sequence of characters that flows to and from a device

- Streams can either be text or binary
- Streams are buffered by default
- Streams are manipulated with a `FILE*` pointer

4.2 Stream Functions

- `FILE* fopen(filename, mode);`
Opens a stream to the specified file in the specified access mode
- `int fclose(stream);`
Closes the specified stream and file
- `int fprintf(stream, format, ...);`
Writes a formatted C string, similar to `printf(...)`
- `int fscanf(stream, format, ...);`
Reads data and stores data matching the format string
- `size_t fwrite(ptr, size, count, stream);`
Writes an array of `count` elements of `size` bytes from `ptr` to `stream`
- `size_t fread(ptr, size, count, stream);`
Reads an array of `count` elements of `size` bytes from `stream` to `ptr`

4.3 Stream Error Checking & Handling

- `int ferror(stream);`
Checks if the error indicator associated with the specified stream is set
- `int clearerr(stream);`
Reset error and EOF indicators for the specified stream
- `void perror(message);`
Prints `message` followed by an error message related to `errno` to `stderr`

4.4 Stream Buffering

By default, data written by `fwrite()` is copied into a buffer inside the process' address space

- At some point, the buffer will be drained into the destination
 - When the buffer size is exceeded
 - When `fflush()` is explicitly called
 - When `fclose()` is called
 - When the process exits gracefully
- Advantages of stream buffering
 - Buffering improves performance by reducing disk accesses
 - Buffering improves abstraction and convenience
- Disadvantages of stream buffering
 - Buffering reduces reliability since the buffer needs to be flushed
 - Buffering reduces performance by requiring extra copies from the buffer to the file

4.5 Portable Operating System Interface (POSIX)

POSIX is a set of lower-level file access APIs

- `int open(filename, mode);`
Opens a stream to the specified file in the specified access mode
- `int close(stream);`
Closes the specified stream and file
- `ssize_t read(int fd, void* buf, size_t count);`
Advances forward in the file by number of bytes read and returns the number of bytes read

4.6 Directories

A directory is a special file that stores the names and locations of the related files/directories

- This includes itself (`.`), its parent directory (`..`), and all of its children
- The directory is accessible via POSIX

5 Programming in C++

5.1 Pointers

Pointers are special variables that store addresses

- Pointers are initialized via `int* x_ptr = &x;`
- Modifying the pointer does not modify what it points to
- Pointers in C++ work the same as in C

5.2 References

References are aliases for other variables

- References are initialized via `int& x_ref = &x;`
- Modifying the reference modifies the aliased variable
- `const` references are used for complex structs/object instances to reduce memory use

5.3 `const`

The `const` keyword indicates to the compiler that the associated variable cannot be mutated

- `const` next to the pointer name indicates that the pointer value cannot be mutated
 - i.e. `int* const ptr`
- `const` next to the data type indicates that the object being referenced by the pointer cannot be mutated
 - i.e. `const int* ptr`

5.4 Classes

- Objects can be declared as `const`
 - Once a `const` object has been constructed, its member variables cannot be changed
 - A `const` object can only invoke member functions that are labeled `const`
 - i.e. `const Name identifier`
- Object member functions can be declared as `const`
 - A `const` member function cannot modify the object it was called on
 - Member functions that do not modify the object should be marked `const`
 - i.e. `retType methodName() const { // body statements; }`

5.5 Class Definitions

```

1  class Name {
2      public:
3      // public member definitions and declarations
4
5      private:
6      // private member definitions and declarations
7  };
8
9  // friend function definitions that use the class

```

- Class definitions are contained in a `.h` file
- Members can be functions or variables
- In-line setter and getter methods are also declared in the header file

5.6 Class Declarations

```

1  retType Name::MethodName(type1 param1, ..., typeN paramN) {
2      // body statements
3  }

```

- Class declarations are contained in a `.cc` file

5.7 Constructors

A constructor (ctor) initializes a newly-instantiated object

- A class can have multiple constructors that differ in parameters
- A constructor must be invoked when creating a new instance of an object
- Classes often have a default constructor, an explicit constructor, and a copy constructor
 - The default copy constructor performs a shallow copy of all the fields

5.8 Destructors

A destructor (dtor) frees any dynamic storage or other resources owned by the object

- The destructor is automatically invoked when a class instance is deleted or goes out of scope

5.9 Rule of Three

If a destructor, copy constructor, or assignment operator is defined, then all three must be defined

5.10 Access Modifiers

Access modifiers control the visibility of fields, methods, and constructors in a class

- `public` members are accessible to all parts of the program
- `private` members are accessible to other member functions of the class
- `protected` members are accessible to member functions of the class and any sub-classes

`struct` members default to `public` and `class` members default to `private`

5.11 Non-Member Functions

Non-member functions are regular functions that make use of some class

- Non-member functions are called as regular functions instead of as a member of a class object instance
- Non-member functions do not have access to the class' private members
- Non-member functions are declared in the class' header file, but outside of the class definition

5.12 friend Non-Member Functions

A class can give a non-member function access to its non-public members by declaring it as a `friend` within its definition

- `friend` non-member functions have the same access privileges as a member function
- `friend` functions are usually unnecessary if the class includes the appropriate public getter functions

5.13 Namespaces

A namespace defines a scope to the identifiers inside it

- If a namespace already exists, then re-defining the namespace adds to the existing namespace

5.14 new/delete

To allocate memory on the heap in C++, use the `new` keyword instead of `malloc()`. To deallocate memory on the heap in C++, use the `delete` keyword instead of `free()`

- `new` never returns a null pointer
- i.e. `int* type_ptr = new int;` allocates a new primitive integer type and `delete type_ptr;` deallocates the type
- i.e. `Point* object_ptr = new Point();` allocates a new `Point` object and `delete object_ptr;` deallocates the object
- i.e. `int* arr_ptr = new int[size];` allocates a new primitive integer array and `delete[] arr_ptr;`

5.15 Templates

Templates enable functions or classes to accept generic types as parameters

```

1  template <typename T>
2  int compare(const T &value1, const T &value2) {
3      ...
4  }
```

- Templated functions should be declared in the header file alongside their definitions

6 C++ Standard Template Library

6.1 Type Inference

Types can be inferred using the `auto` keyword

```

1   type Function();
2
3   int main(int argc, char* argv[]) {
4       auto var = Function();
5   }
```

- The expression using `auto` must contain explicit initialization with defined return types

6.2 Containers

A container is an object that stores a collection of other objects in memory

- Containers store by value instead of by reference
- When an object is inserted, the container makes a copy
- Sequence containers index their elements numerically
 - i.e. `vector`, `deque`, `list`
- Associative containers index their elements by key
 - i.e. `set`, `map`, `multiset`

6.3 Lists

A list is a generic doubly-linked list

- Does not support random access
- Inserting and deleting elements is $O(1)$ time

6.4 Vectors

A vector is a generic dynamically resizable contiguous array

- Random access is $O(1)$ time
- Inserting and deleting elements from the end is amortized $O(1)$ time
- Inserting and deleting elements from the start or middle is $O(n)$ time

6.5 Maps

A map is a key/value table implemented as a search tree

- Elements are type `pair<key_type, value_type>` and are stored in sorted order
- `key_type` must support less-than operator
- Lookup and insertion is $O(\log n)$ time

6.6 Iterators

Each container class has an associated iterator class used to iterate through elements of the container

- Iterators range from `begin()` to `end()`, where `end()` is one past the last container element

7 Smart Pointers

7.1 Smart Pointers

A smart pointer is an object that stores a pointer to a heap-allocated object

- Smart pointers look and behave like a regular pointer
- Smart pointers will destroy and delete the associated object at the right time

7.2 `shared_ptr`

`shared_ptr` is a smart pointer that tracks the number of references to a piece of data, and only deallocates when no smart pointers are managing that data

- Constructors create the counter
- Copy constructors and assignment operators increment the counter
- Destructors decrement the counter

7.3 `unique_ptr`

`unique_ptr` is a smart pointer that is the sole owner of a pointer to a piece of data

- `unique_ptr` can release ownership of a pointer via `release()`
 - `x.release()` returns the pointer stored in `x` and replaces it with `nullptr`
- A new `unique_ptr` can accept ownership of a pointer via initialization
 - `y(x.release())` constructs a new `unique_ptr` with the pointer stored in `x`
- An existing `unique_ptr` can accept ownership of a pointer via `reset()`
 - `y.reset(x.release())` deletes the previous pointer and its associated data, and replaces it with the pointer stored in `x`
- `unique_ptr` can transfer ownership of a pointer via `move()`
 - `y = move(x)` interface is equivalent to `y.reset(x.release())`

7.4 `weak_ptr`

`weak_ptr` is similar to `shared_ptr` but does not affect the reference count

- `weak_ptr` cannot be directly dereferenced

8 Inheritance

8.1 Inheritance

A derived/child class inherits all the methods and properties from a base/parent class

```
1  #include "BaseClass.h"
2
3  class DerivedClass : public BaseClass {
4      ...
5  }
```

- Constructors, destructors, copy constructor, and assignment operator are never inherited

8.2 Virtual Functions

A virtual function is a member function of a base class which may be redefined by derived classes

- It is best practice for derived functions to declare member functions as `virtual` where appropriate

8.3 Overriding Functions

Virtual functions can be overridden using the `override` keyword

- Declares to the compiler that this method should be overriding an inherited virtual function
- Prevents overloading vs. overriding bugs
- A member function cannot be declared as both `virtual` and `override`

8.4 Pure Virtual Functions

A pure virtual function is a member function of a base class which is only implemented in derived classes

- A pure virtual function is defined via `virtual retType MethodName() = 0;`

8.5 Abstract Classes

A class containing any pure virtual function is abstract

- Abstract classes cannot be instantiated
- Abstract classes must be extended and overridden in order to be used
- An interface is an abstract class containing only pure virtual functions

9 Casts

9.1 Static Casts

`static_cast<to_type>(expression)` performs well-defined conversions

- Can convert pointers between classes of related type
 - i.e. casting between `void*` and `T*`
- Can convert related references
 - i.e. casting between `float` and `int`
- Static casts are checked at compile time

9.2 Dynamic Casts

`dynamic_cast<to_type>` performs related conversions

- Can convert pointers between classes of related type
 - i.e. casting between `BaseClass*` and `DerivedClass*`
- Can convert references between classes of related type
 - i.e. casting between `BaseClass` and `DerivedClass`
- Dynamic casts are checked at both compile time and run time

9.3 Constant Casts

`const_cast<to_type>` adds or removes const-ness

- Can convert between `const` and non-`const`

9.4 Reinterpret Casts

`reinterpret_cast<to_type>` performs conversions between incompatible types

- Involves the low-level reinterpretation of the bit pattern
 - i.e. storing a pointer in an `int`